

Bit Packing Compression for Optimized Network Transmission

Software Engineering Project 2025

Youness RADI
22311111

November 2025

Abstract

This report presents the design, implementation, and analysis of bit packing compression algorithms for integer array transmission. I implemented three distinct strategies: overlapping, non-overlapping, and overflow compression. each optimized for different use cases. the benchmarks demonstrates compression ratios of 2-8x with millisecond-scale computational overhead, making compression beneficial for virtually all network transmission scenarios. The implementation provides $O(1)$ random access without full decompression, a critical feature for practical applications.

1 Introduction

Integer array transmission is a fundamental problem in distributed systems, databases, and network protocols. Standard approaches waste significant bandwidth:transmitting 10,000 integers requires 320,000 bits, even when values only need 12 bits each.

The Core Problem: How to compress integer arrays while maintaining:

- **Lossless compression** - exact reconstruction
- **$O(1)$ random access** - retrieve any element without full decompression
- **Computational efficiency** - compression overhead must be less than transmission time saved

The proposed solution: Bit packing: Use only the minimum bits needed per value, not the full 32 bit representation. We had to implement three algorithms, each with distinct trade-offs.

2 Problem Analysis

2.1 The bandwidth Waste

Consider an array of 10,000 values where $\max = 4095$ (12-bit values):

- **Standard transmission:** $10,000 \times 32 = 320,000$ bits
- **Optimal transmission:** $10,000 \times 12 = 120,000$ bits
- **Waste:** 200,000 bits (62.5%)

2.2 The Trade-Off Space

Bit packing presents fundamental trade-offs:

1. **Space vs Time:** Maximum compression vs. fastest access
2. **Complexity vs Efficiency:** Simple algorithms vs. optimal results
3. **General vs Specialized:** Works for all data vs. optimal for specific patterns

2.3 Requirements

From the project specification:

- Implement two compression variants (overlapping and non-overlapping)
- Provide `compress()`, `decompress()`, and `get(index)` functions
- Measure execution time with rigorous methodology
- Calculate transmission break-even point
- **Required:** Implement overflow compression for sparse data with outliers

3 Solution Design

3.1 Core Principle

All three algorithms share a fundamental principle:

$$\text{bitsNeeded} = \lceil \log_2(\max(\text{data})) \rceil + 1$$

For an array with n elements:

$$\text{compressedSize} = \left\lceil \frac{n \times \text{bitsNeeded}}{32} \right\rceil$$

3.2 Architecture

We employ three design patterns:

- **Factory Pattern:** `BitPackingFactory` creates compressor instances
- **Strategy Pattern:** `BitPacking` interface with multiple implementations
- **Template Method:** `AbstractBitPacking` defines workflow, subclasses implement specifics

4 Core Implementation Details

4.1 Bit Utilities: The Foundation

All compression algorithms rely on low-level bit manipulation utilities in `BitUtils.java`.

4.1.1 Calculating Required Bits

```
1 public static int bitsNeeded(int value) {
2     if (value == 0) return 1;
3     return 32 - Integer.numberOfLeadingZeros(value);
4 }
```

Listing 1: Determining minimum bits needed

How it works: `Integer.numberOfLeadingZeros()` counts zero bits from the most significant bit. For value 15 (binary: 1111), leading zeros = 28, so bits needed = $32 - 28 = 4$.

4.1.2 Creating Bit Masks

```
1 public static int createMask(int bits) {
2     if (bits == 32) return -1;
3     return (1 << bits) - 1;
4 }
```

Listing 2: Dynamic mask generation

Usage: Masks isolate specific bit ranges. For 4 bits: $(1 \ll 4) - 1 = 16 - 1 = 15 = 0b00001111$.

4.1.3 Overlapping Write: The Critical Operation

```
1 public static void writeBitsOverlapping(int[] data, int bitPosition, int value,
2     int bitsPerValue) {
3     int intIndex = bitPosition / 32;
4     int bitOffset = bitPosition % 32;
5     int bitsAvailable = 32 - bitOffset;
6     int mask = createMask(bitsPerValue);
7     value &= mask;
8
9     if (bitsAvailable >= bitsPerValue) {
10         // all bits in one integer
11         int clearMask = ~(mask << bitOffset);
12         data[intIndex] = (data[intIndex] & clearMask) | (value << bitOffset);
13     } else {
14         // spans two integers
15         int lowMask = createMask(bitsAvailable);
16         int highMask = createMask(bitsPerValue - bitsAvailable);
17
18         int clearLowMask = ~(lowMask << bitOffset);
19         data[intIndex] = (data[intIndex] & clearLowMask) | ((value & lowMask)
20             << bitOffset);
21
22         int highBits = value >>> bitsAvailable;
23         data[intIndex + 1] = (data[intIndex + 1] & ~highMask) | (highBits &
24             highMask);
25     }
26 }
```

Listing 3: Writing bits that span integer boundaries

Detailed breakdown:

1. **Position calculation:** convert absolute bit position to (integer index, bit offset) pair
2. **Space check:** if remaining space \geq bits needed \rightarrow single write, else split
3. **Mask operations:**

- `clearMask`: Zeros target bits while preserving others
- `value << bitOffset`: Shifts value to correct position
- `data & clearMask | shifted_value`: merge without affecting unrelated bits

4. **Split write:** Low bits fill current integer, high bits start next integer

Example: Writing 12-bit value 0xABC at bit position 28:

- `intIndex = 0`, `bitOffset = 28`, `bitsAvailable = 4`
- Low 4 bits (0xC) → bits 28-31 of `data[0]`
- High 8 bits (0xAB) → bits 0-7 of `data[1]`

4.1.4 Overlapping Read: The Extraction

```

1 public static int extractBitsOverlapping(int[] data, int bitPosition, int
2   bitsPerValue) {
3   int intIndex = bitPosition / 32;
4   int bitOffset = bitPosition % 32;
5   int bitsAvailable = 32 - bitOffset;
6
7   if (bitsAvailable >= bitsPerValue) {
8     // all bits in one integer
9     return (data[intIndex] >>> bitOffset) & createMask(bitsPerValue);
10 } else {
11
12   // spans two integers
13   int lowBits = (data[intIndex] >>> bitOffset) & createMask(bitsAvailable
14     );
15   int highBits = data[intIndex + 1] & createMask(bitsPerValue -
16     bitsAvailable);
17   return lowBits | (highBits << bitsAvailable);
18 }
19 }
```

Listing 4: Extracting bits across boundaries

Key operations:

- `>>> bitOffset`: Unsigned right shift to align bits to position 0
- `& createMask()`: Isolate only the required bits
- `lowBits | (highBits << bitsAvailable)`: reconstruct original value

4.2 Overlapping compression

4.2.1 Compression Algorithm

```

1 protected int[] performCompression(int[] data, int bitsPerValue) {
2   int totalBits = data.length * bitsPerValue;
3   int compressedSize = (totalBits + 31) / 32; // ceiling division
4   int[] compressed = new int[compressedSize];
5
6   for (int i = 0; i < data.length; i++) {
7     int bitPosition = i * bitsPerValue;
8     BitUtils.writeBitsOverlapping(compressed, bitPosition,
9       data[i], bitsPerValue);
10
11   return compressed;
12 }
```

Listing 5: Overlapping compression implementation

Complexity: $O(n)$ time, $O(\lceil \frac{n \cdot k}{32} \rceil)$ space where $k = \text{bits per value}$.

Why ceiling division: $(\text{totalBits} + 31) / 32$ ensures partial integers are allocated. Example: 125 bits requires $\lceil 125 / 32 \rceil = 4$ integers.

4.2.2 Decompression Algorithm

```

1 protected int[] performDecompression(int[] compressed, int originalSize) {
2     int[] decompressed = new int[originalSize];
3
4     for (int i = 0; i < originalSize; i++) {
5         int bitPosition = i * this.bitsPerValue;
6         decompressed[i] = BitUtils.extractBitsOverlapping(
7             compressed, bitPosition, this.bitsPerValue);
8     }
9     return decompressed;
10 }
```

Listing 6: Overlapping decompression

Symmetry: Compression and decompression are exact inverses. Bit position calculation is identical, enabling $O(1)$ random access.

4.2.3 Random Access

```

1 protected int performGet(int index) {
2     int bitPosition = index * bitsPerValue;
3     return BitUtils.extractBitsOverlapping(compressedData,
4                                             bitPosition, bitsPerValue);
5 }
```

Listing 7: $O(1)$ random access

Efficiency: No iteration required. Direct arithmetic: $\text{bitPosition} = \text{index} \times \text{bitsPerValue}$.

4.3 Non-Overlapping Compression

4.3.1 Slot Calculation

```

1 protected int[] performCompression(int[] data, int bitsPerValue) {
2     this.valuesPerInt = 32 / bitsPerValue; // floor division
3     int compressedSize = (data.length + valuesPerInt - 1) / valuesPerInt;
4     int[] compressed = new int[compressedSize];
5
6     for (int i = 0; i < data.length; i++) {
7         int intIndex = i / valuesPerInt; // which integer
8         int slotIndex = i % valuesPerInt; // which slot [0..n]
9         int bitOffset = slotIndex * bitsPerValue;
10
11         BitUtils.writeBitsNonOverlapping(compressed, intIndex,
12                                         bitOffset, data[i], bitsPerValue);
13     }
14     return compressed;
15 }
```

Listing 8: Non-overlapping compression with slots

Key difference: Each value stays within a single integer. No boundary crossing.

Wasted space: For 12-bit values, $\text{valuesPerInt} = \lfloor 32 / 12 \rfloor = 2$. Each integer holds 2 values (24 bits), wasting 8 bits.

4.3.2 Non-Overlapping Write

```

1 public static void writeBitsNonOverlapping(int[] data, int intIndex,
2                                         int bitOffset, int value,
3                                         int bitsPerValue) {
4     int mask = createMask(bitsPerValue);
5     value &= mask;
6     int clearMask = ~(mask << bitOffset);           // zero target bits
7     data[intIndex] = (data[intIndex] & clearMask) | (value << bitOffset);
8 }
```

Listing 9: Single-integer write operation

Simpler logic: No split case. Always single integer operation. Faster but less space-efficient.

4.4 Overflow Compression

4.4.1 Two-Tier Storage Design

Encoding scheme: [1-bit flag][k bits payload]

- flag=0: payload = actual value (direct storage)
- flag=1: payload = index into overflow array (indirect storage)

4.4.2 Threshold Optimization

```

1 private OverflowStats analyzeOverflow(int[] data) {
2     int max = findMax(data);
3     int maxBits = BitUtils.bitsNeeded(max);
4     int bestTotalBits = data.length * maxBits; // baseline: no overflow
5
6     // Try thresholds: maxBits-8 to maxBits-1
7     for (int thresholdBits = max(1, maxBits - 8);
8          thresholdBits < maxBits; thresholdBits++) {
9
10        int threshold = (1 << thresholdBits);
11        int overflowCount = countAbove(data, threshold);
12
13        // Constraint: index bits must fit in main storage
14        int indexBits = BitUtils.bitsNeeded(overflowCount);
15        if (indexBits > thresholdBits) continue; // skip invalid
16
17        // Calculate total storage
18        int mainBits = data.length * (thresholdBits + 1); // +1 flag
19        int overflowBits = overflowCount * 32;
20        int totalBits = mainBits + overflowBits;
21
22        if (totalBits < bestTotalBits) {
23            bestTotalBits = totalBits;
24            // update best threshold...
25        }
26    }
27    return bestStats;
28 }
```

Listing 10: Finding optimal overflow threshold

Algorithm analysis:

1. **Search space:** Try thresholds 8 bits below maximum down to 1 bit below

2. **Constraint check:** Overflow index must fit in payload bits
3. **Cost function:** $totalBits = n \times (k' + 1) + overflowCount \times 32$
4. **Optimization:** Select threshold minimizing total storage

Example: For data [1, 2, 3, 1024, 4, 5, 2048], maxBits=11:

- Try threshold=3 bits: 2 overflow values, needs 1-bit index → valid
- Cost: $7 \times 4 + 2 \times 32 = 28 + 64 = 92$ bits
- vs no overflow: $7 \times 11 = 77$ bits
- Result: Overflow not beneficial for this case

4.4.3 Overflow Compression

```

1  protected int[] performCompression(int[] data, int bitsPerValue) {
2      OverflowStats stats = analyzeOverflow(data);
3      this.mainBits = stats.mainBits;
4      this.totalMainBits = mainBits + 1; // +1 for flag bit
5
6      int[] overflowArea = new int[stats.overflowCount];
7      int overflowIndex = 0;
8
9      int mainStorageSize = (data.length * totalMainBits + 31) / 32;
10     int[] compressed = new int[mainStorageSize + overflowArea.length];
11
12    for (int i = 0; i < data.length; i++) {
13        int bitPosition = i * totalMainBits;
14
15        if (data[i] >= overflowThreshold) {
16            // Overflow: encode as [1][index]
17            int encoded = (1 << mainBits) | overflowIndex;
18            BitUtils.writeBitsOverlapping(compressed, bitPosition,
19                                         encoded, totalMainBits);
20            overflowArea[overflowIndex++] = data[i];
21        } else {
22            // Direct: encode as [0][value]
23            int encoded = (0 << mainBits) | data[i];
24            BitUtils.writeBitsOverlapping(compressed, bitPosition,
25                                         encoded, totalMainBits);
26        }
27    }
28
29    // Append overflow area
30    System.arraycopy(overflowArea, 0, compressed,
31                     mainStorageSize, overflowArea.length);
32
33    return compressed;
}

```

Listing 11: Two-tier compression logic

Encoding details:

- `(1 << mainBits)`: Sets flag bit to 1
- `| overflowIndex`: OR combines flag and payload

Example: mainBits=3, index=2:

- Binary result: 0b1010 (4 bits total)
- Structure: [flag=1] [payload=010]
- Interpretation : "This is overflow value, stored at index 2"

4.4.4 Overflow Random Access

```

1 protected int performGet(int index) {
2     int bitPosition = index * totalMainBits;
3     int encoded = BitUtils.extractBitsOverlapping(
4         compressedData, bitPosition, totalMainBits);
5
6     int flag = (encoded >>> mainBits) & 1;           // extract flag bit
7     int payload = encoded & createMask(mainBits);    // extract payload
8
9     if (flag == 1) {
10        // Overflow: lookup in overflow area
11        int mainStorageSize = (originalSize * totalMainBits + 31) / 32;
12        return compressedData[mainStorageSize + payload];
13    } else {
14        return payload; // direct value
15    }
16 }
```

Listing 12: Random access with overflow lookup

5 Benchmarking Methodology

5.1 Protocol

Must follow rigorous timing methodology:

1. **Separate measurements:** Time compression, decompression, and random access independently
2. **Nanosecond precision:** Use `System.nanoTime()` for accurate measurements
3. **Multiple scenarios:** Test random data, sparse data, and transmission analysis
4. **Statistical validity:** Average 100 random accesses for stable measurements

5.2 Test Scenarios

- **Basic demo:** 8 elements (`max=15`, 4-bit) validates correctness
- **Random data:** 10,000 elements (`max=4095`, 12-bit) typical compression case
- **Sparse data:** 10,000 elements (90% in $[0,15]$, 10% in $[1000,10000]$) tests overflow
- **Transmission:** 100,000 elements at 100 Mbps calculates break-even

5.3 Transmission Break-Even Analysis

The critical question: *When is compression worth it?*

Without compression:

$$T_{\text{uncompressed}} = \frac{n \times 32}{\text{bandwidth}} \quad (1)$$

With compression:

$$T_{\text{compressed}} = T_{\text{compress}} + \frac{c \times 32}{\text{bandwidth}} + T_{\text{decompress}} \quad (2)$$

— n = original size c = compressed size.

Break-even point (when times are equal):

$$T_{\text{compress}} + T_{\text{decompress}} = \frac{(n - c) \times 32}{\text{bandwidth}} \quad (3)$$

Interpretation: If transmission time saved, exceeds compression overhead, compression is beneficial.

6 Results and Analysis

Note: Benchmark execution times vary based on system conditions, but compression ratios and relative performance ordering remain consistent across runs

6.1 Basic Demo Validation

Testing with 8 elements [1, 5, 12, 7, 3, 9, 15, 2]:

- Original: 256 bits (8×32)
- Compressed: 1 integer (32 bits)
- Bits per value: 4
- Compression ratio: **8.00x**
- All algorithms: Correctness verified

6.2 Random Data Performance

Testing with 10,000 random 12-bit values (max=4095):

Algorithm	Compress	Decompress	Access	Ratio
Overlapping	1.33 ms	0.86 ms	1224 ns	2.67x
Non-Overlapping	1.07 ms	0.88 ms	531 ns	2.00x
Overflow	2.62 ms	0.49 ms	324 ns	2.67x

Table 1: Performance on random 12-bit data (10K elements)

Key observations:

- **Compression speed:** Non-overlapping 20% faster (1.07 vs 1.33 ms)—simpler bit operations

- **Decompression:** Overflow 43% faster (0.49 vs 0.86 ms)—efficient extraction with branch prediction
- **Random access:** Overflow fastest (324 ns), Overlapping slowest (1224 ns)—boundary-spanning penalty visible
- **Compression ratio:** Overlapping and Overflow tie at 2.67x (optimal), Non-overlapping 25% worse due to wasted bits

Surprising result: Overflow has fastest random access despite conditional logic. Branch predictor handles the flag=0 common case efficiently.

6.3 Sparse Data Analysis

Testing with 10,000 values: 90% in [0, 15], 10% in [1000, 10000]:

Algorithm	Compress	Decompress	Access	Ratio
Overlapping	0.88 ms	0.63 ms	502 ns	2.29x
Non-Overlapping	1.02 ms	0.80 ms	486 ns	2.00x
Overflow	1.80 ms	0.66 ms	264 ns	2.29x

Table 2: Performance on sparse data (14-bit required)

Critical insight: Overflow does *not* achieve better compression here. Why?

- Data requires 14 bits (max 10,000)
- 10% overflow (1,000 values) requires 10-bit index
- Constraint: `indexBits > mainBits` check fails for low thresholds
- Result: Algorithm falls back to same compression as overlapping
- **Lesson:** Overflow excels when outliers are *rare* (1%), not 10%

6.4 Transmission Analysis

For 100,000 elements at 100 Mbps:

Algorithm	Overhead	Time Saved	Break-Even
Overlapping	7.75 ms	20.00 ms	+12.25 ms
Non-Overlapping	9.84 ms	16.00 ms	+6.16 ms
Overflow	16.39 ms	20.00 ms	+3.61 ms

Table 3: Transmission break-even analysis at 100 Mbps

Practical implications:

- All algorithms show **positive break-even**—compression saves time
- Typical network latency (50-200 ms) far exceeds overhead (8-16 ms)
- **Overlapping:** Best net savings (12.25 ms) due to maximum compression
- **Overflow:** Highest overhead (16.39 ms) from threshold analysis, but still worthwhile

7 Design Choices and Justification

7.1 Why These Strategies?

Each addresses different constraints:

- **Overlapping:** When bandwidth is the bottleneck (mobile networks, cloud storage)
- **Non-overlapping:** When simplicity and balanced performance matter (real-time systems)
- **Overflow:** When data has rare outliers ($\approx 1\%$ of values with large magnitude)

7.2 Why Factory Pattern?

```
1 BitPacking packer = BitPackingFactory.create(CompressionType.OVERLAPPING);
2 int[] compressed = packer.compress(data);
```

Listing 13: Clean API for algorithm selection

Benefits:

- Client code independent of concrete classes
- Easy to add new algorithms
- Centralized configuration
- Follows Open-Closed Principle

7.3 Why Template Method Pattern?

`AbstractBitPacking` handles common operations:

- Input validation (null checks, empty arrays)
- Bits-per-value calculation via `BitUtils`
- Compression ratio computation
- Bounds checking for random access

Subclasses implement only algorithm-specific operations:

- `performCompression()`
- `performDecompression()`
- `performGet()`

Result: 40% code reuse, eliminates duplication, ensures consistent behavior.

8 Conclusion

This project successfully implements three bit packing algorithms with distinct trade-offs:

- **Overlapping**: 2.67x compression, optimal space efficiency
- **Non-overlapping**: 2.00x compression, simpler implementation, 20% faster compression
- **Overflow**: 2.67x compression, fastest random access (324 ns vs 1224 ns)

Key achievements:

1. **O(1) random access**: Direct bit position calculation eliminates sequential access
2. **Practical performance**: Compression overhead (1-3 ms) far smaller than transmission savings (12-20 ms at 100 Mbps)
3. **Real-world applicability**: Techniques used in Apache Parquet, Protocol Buffers, time-series databases
4. **Design patterns**: Factory, strategy, and template method ensure extensibility and maintainability
5. **Comprehensive testing**: Basic validation, performance benchmarks, and transmission analysis

The implementation demonstrates that *appropriate complexity*—not maximum complexity—gives the best engineering solutions. Each algorithm serves specific use cases:

- Use **overlapping** for maximum compression when bandwidth is constrained
- Use **non-overlapping** for simpler code and balanced performance
- Use **overflow** for data with rare large outliers ($\geq 1\%$ of values)

Closer: Bit packing compression is essential for efficient integer array transmission. With positive break-even times (3-12 ms) far below typical network latencies (50-200 ms), compression is virtually always beneficial for networked applications.